

Linux Internals

Linux Inter Process Communication

Contents

- Module Coverage
 - What is IPC?
 - Different IPC mechanisms
 - Shared Memory
 - Mapped Memory (mmap)
 - Pipe and FIFO
 - Semaphores
 - Sockets

What is IPC?

- Interprocess communication (IPC) is the transfer of data among processes
- This transfer of data can be between, related processes (parent-child relationship), unrelated processes, and processes on different machines
- Example:
 - We use the command "`ls | lpr`" to print the filenames in a directory
 - The `ls` process writes data into the pipe, and the `lpr` process reads data from the pipe.

Different IPC mechanisms

- The different IPC mechanism available in Linux are:
 - Pipes
 - Pipes permit sequential communication from one process to a related process
 - FIFOs
 - FIFOs are similar to pipes, except that unrelated processes can communicate because the pipe is given a name in the filesystem
 - Message Queues
 - Message Queues can be used by processes to read and write messages
 - Shared memory
 - Shared memory allows processes to communicate by reading and writing to a specified memory location.

Different IPC mechanisms (Contd..)

- Mapped memory
 - Mapped memory is similar to shared memory, except that it is associated with a file in the filesystem
 - Contents of a file are mapped to memory
- Semaphores
 - Semaphores are used to synchronize access to shared resources
- Sockets
 - Sockets are used for communication between processes on same or different computers

Pipes

- A pipe is a communication device that permits unidirectional communication
- Typically, a pipe is used to communicate between two threads in a single process or between parent and child processes (i.e. related processes)
- In a shell, the symbol `|` presents a pipe. For example, this shell command causes the shell to produce two child processes, one for `ls` and one for `less`. The two child processes communicate using a pipe.
 - `% ls | less`
- A pipe's data capacity is limited. The pipe automatically synchronizes the two processes

Creating Pipes

- To create a pipe, invoke the pipe system call.
 - Supply an integer array of size 2
 - The call to pipe stores the reading file descriptor in array position 0 and the writing file descriptor in position 1
- For example, consider this code:

```
int pipe_fds[2];  
int read_fd;  
int write_fd;  
pipe (pipe_fds);  
read_fd = pipe_fds[0];  
write_fd = pipe_fds[1];
```

Communication Between Parent and Child Processes

- Parent creates pipe
- A fork spawns a child process
- The child inherits the pipe file descriptors
- The parent writes a data to the pipe
- and the child reads it out

FIFOs

- A *first-in, first-out (FIFO)* file is a pipe that has a name in the filesystem.
- Any process can open or close the FIFO; the processes on either end of the pipe need not be related to each other
- FIFOs are also called *named pipes*
- You can make a FIFO using the `mkfifo` command or system call

```
% mkfifo /tmp/fifo
% ls -l /tmp/fifo
prw-rw-rw- 1 samuel users 0 Jan 16 14:04 /tmp/fifo
```
- The first character of the output from `ls` is `p`, indicating that this file is actually a FIFO (named pipe)

Creating FIFO

- Create a FIFO programmatically using the `mkfifo` function.
- The first argument is the path at which to create the FIFO; the second parameter specifies the pipe's owner, group, and world permissions
- If the pipe cannot be created (for instance, if a file with that name already exists), `mkfifo` returns `-1`
- Include `<sys/types.h>` and `<sys/stat.h>` if you call `mkfifo`

Accessing FIFO

- A FIFO is accessed just like an ordinary file
 - To communicate through a FIFO, one program must open it for writing, and another program must open it for reading
- Either low-level I/O functions (open, write, read, close, and so on, or C library I/O functions (fopen, fprintf, fscanf, fclose, and so on) may be used
 - For example to write a buffer of data to a FIFO using low-level I/O

```
int fd = open (fifo_path, O_WRONLY);
write (fd, data, data_length);
close (fd);
```
 - To read a string from the FIFO using C library I/O functions

```
FILE* fifo = fopen (fifo_path, "r");
fscanf (fifo, "%s", buffer);
fclose (fifo);
```

Shared Memory

- Shared memory allows two or more processes to access the same memory.
- Access to the shared memory is as fast as accessing a process's nonshared memory, and it does not require a system call or entry to the kernel.
- All processes share the same piece of memory.
- The kernel does not synchronize accesses to shared memory, we must provide our own synchronization mechanisms. A common strategy is to use semaphores.

Shared Memory Issues

- Shared memory segments allow fast communication among any number of processes
- Each process can both read and write, but must follow a protocol to avoid race conditions
- Processes must make arrangements to use the same key to use a shared segment

The memory model

- To use a shared memory segment, one process must allocate the segment
- Then each process desiring to access the segment must attach the segment
- After finishing its use of the segment, each process detaches the segment.
- At some point, one process must deallocate the segment

Shared memory Data Structure

```
struct shmid_ds
{
    struct ipc_perm shm_perm;    /* operation perms */
    int  shm_segsz;      /* size of segment (bytes) */
    time_t shm_atime;    /* last attach time */
    time_t shm_dtime;    /* last detach time */
    time_t shm_ctime;    /* last change time */
    unsigned short shm_cpid; /* pid of creator */
    unsigned short shm_lpid; /* pid of last operator */
    short shm_nattch; /* no. of current attaches */
    /*
    .....
    */
};
```

Shared Memory Allocation

- A process allocates a shared memory segment using `shmget(key,size,options)` system call
- Its first parameter is an integer key that specifies which segment to create
- Its second parameter specifies the number of bytes in the segment.
- The third parameter is the bitwise or of flag values that specify options to `shmget`
 - `IPC_CREAT`
 - `IPC_EXCL`
 - Mode flags

shmget() system call

```
int shmget ( key_t key, int size, int shmflg );
```

returns: shared memory segment identifier on success

Arg1 : keyvalue

Arg2 : size

Arg3 :

IPC_CREAT

Create the segment if it doesn't already exist in the kernel.

IPC_EXCL

When used with IPC_CREAT, fail if segment already exists.

Shared Memory Allocation (Contd..)

- For example

```
int segment_id = shmget (shm_key, getpagesize (),  
    IPC_CREAT | S_IRUSR | S_IWUSR );
```

- If the call succeeds, shmget returns a segment identifier
- If the shared memory segment already exists, the access permissions are verified and a check is made to ensure that the segment is not marked for destruction.

Attachment and Detachment

- To make the shared memory segment available, a process must use `shmat(shmid,address,flags)`
 - Pass it the shared memory segment identifier `shmid` returned by `shmget`
 - The second argument is a pointer that specifies where in your process's address space you want to map the shared memory; if you specify `NULL`, Linux will choose an available address
 - The third argument is a flag, which can include the following:
 - `SHM_RND` indicates that the address specified for the second parameter should be rounded down to a multiple of the page size. If you don't specify this flag, you must page-align the second argument to `shmat` yourself
 - `SHM_RDONLY` indicates that the segment will be only read, not written

Attachment of a Segment

- SYSTEM CALL: `shmat()`;
PROTOTYPE: `int shmat (int shmid, char *shmaddr, int shmflg);`
- RETURNS: address at which segment was attached to the process, or `-1` on error
- `char *attach_segment(int shmid)`
- {
- `return(shmat(shmid, 0, 0));`
- }

Attachment and Detachment (Contd..)

- If the call succeeds, it returns the address of the attached shared segment
- Children created by calls to fork inherit attached shared segments; they can detach the shared memory segments, if desired.
- When you're finished with a shared memory segment, the segment should be detached using `shmdt(address)`
- Pass it the address returned by `shmat`
- If the segment has been deallocated and this was the last process using it, it is removed. Calls to `exit` and any of the `exec` family automatically detach segments

Controlling and Deallocating Shared Memory

- The `shmctl` call returns information about a shared memory segment and can remove it.
 - The first parameter is a shared memory segment identifier
 - To obtain information about a shared memory segment, pass `IPC_STAT` as the second argument and a pointer to a struct `shmid_ds`
 - To remove a segment, pass `IPC_RMID` as the second argument, and pass `NULL` as the third argument. The segment is removed when the last process that has attached it finally detaches it.
 - Each shared memory segment should be explicitly deallocated using `shmctl` when you're finished with it, to avoid violating the systemwide limit on the total number of shared memory segments. Invoking `exit` and `exec` detaches memory segments but does not deallocate them.

Controlling and Deallocating Shared Memory (contd)

- The `ipcs` command provides information on IPC facilities.
- Use `-m` flag to obtain information about shared memory

- `% ipcs -m`

```
----- Shared Memory Segments -----  
key shmid owner perms bytes nattch  
0x00000000 1627649 user 640 25600 0
```

- Use `ipcrm` command to remove a shm.
- `% ipcrm shm 1627649`

Semaphore

A semaphore is a protected variable and constitutes the classic method for restricting access to shared resources (e.g. storage) in a multiprogramming environment

It supports 2 operations

P(sv)-decrement sv

V(sv)-increment sv

Semaphore definitions

```
#include<sys/sem.h>
```

```
Int semctl(int sem_id,int sem_num,int command...)
```

```
Int semget(key_t key,int num_sems,int sem_flags);
```

```
Int semop(int sem_id,struct sembuf *sem_ops,size_t num_sem_ops);
```

Processes Semaphores

- The calls `semget` and `semctl` allocate and deallocate semaphores
- Invoke `semget` with a key specifying a semaphore set, the number of semaphores in the set, and permission flags as for `shmget`
- The return value is a semaphore set identifier.
- The last process to use a semaphore set must explicitly remove it to ensure that the operating system does not run out of semaphores.
- To do so, invoke `semctl` with the semaphore identifier, the number of semaphores in the set, `IPC_RMID` as the third argument, and any union `semun` value as the fourth argument (which is ignored)

Processes Semaphores (Contd..)

```
/* We must define union semun ourselves. */
union semun {
    int val;
    struct semid_ds *buf;
    unsigned short int *array;
    struct seminfo *__buf;
};
/* Obtain a binary semaphore's ID, allocating if necessary. */
int binary_semaphore_allocation (key_t key, int sem_flags) {
    return semget (key, 1, sem_flags);
}
/* Deallocate a binary semaphore. All users must have finished their
use. Returns -1 on failure. */
int binary_semaphore_deallocate (int semid) {
    union semun ignored_argument;
    return semctl (semid, 1, IPC_RMID, ignored_argument);
}
```

Initializing Semaphores

```
/* We must define union semun ourselves. */
union semun {
    int val;
    struct semid_ds *buf;
    unsigned short int *array;
    struct seminfo *__buf;
};

/* Initialize a binary semaphore with a value of 1. */
int binary_semaphore_initialize (int semid)
{
    union semun argument;
    unsigned short values[1];
    values[0] = 1;
    argument.array = values;
    return semctl (semid, 0, SETALL, argument);
}
```

Wait and Post Operations

- The fields of struct sembuf are listed here:
 - **sem_num**
 - This is the semaphore number in the semaphore set on which the operation is performed
 - **sem_op**
 - This is an integer that specifies the semaphore operation
 - **sem_flg**
 - This is a flag value. Specify IPC_NOWAIT to prevent the operation from blocking; if the operation would have blocked, the call to semop fails instead.

Mapped Memory

- Mapped memory permits different processes to communicate via a shared file
- Although you can think of mapped memory as using a shared memory segment with a name, you should be aware that there are technical differences
- Mapped memory can be used for interprocess communication or as an easy way to access the contents of a file
- Mapped memory forms an association between a file and a process's memory
- You can think of mapped memory as allocating a buffer to hold a file's entire contents, and then reading the file into the buffer and (if the buffer is modified) writing the buffer back out to the file afterward

Mapping an Ordinary File

- `mmap` call is used to map an ordinary file to a process's memory
 - `file_memory = mmap (0, FILE_LENGTH, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);`
- Arguments
 - The first argument is the address at which you would like Linux to map the file into your process's address space; the value `NULL` allows Linux to choose an available start address.
 - The second argument is the length of the map in bytes.
 - The third argument specifies the protection on the mapped address range.

Mapping an Ordinary File (Contd..)

- The fourth argument is a flag value that specifies additional options(MAP_FIXED, MAP_PRIVATE, MAP_SHARED)
 - MAP_FIXED—If you specify this flag, Linux uses the address you request to map the file.This address must be page-aligned.
 - MAP_PRIVATE—Writes to the memory range should not be written back to the attached file, but to a private copy of the file. No other process sees these writes. This mode may not be used with MAP_SHARED.
 - MAP_SHARED—Writes are immediately reflected in the underlying file rather than buffering writes. Use this mode when using mapped memory for IPC.This mode may not be used with MAP_PRIVATE
- The fifth argument is a file descriptor opened to the file to be mapped.
- The last argument is the offset from the beginning of the file from which to start the map

Mapping an Ordinary File (Contd..)

- If `mmap` call succeeds, it returns a pointer to the beginning of the memory. On failure, it returns `MAP_FAILED`.
- When you're finished with a memory mapping, release it by using `munmap`.
 - Pass it the start address and length of the mapped memory region.
 - Linux automatically unmaps mapped regions when a process terminates

MAP_SHARED flag

- Different processes can communicate using memory-mapped regions associated with the same file.
- Specify the MAP_SHARED flag so that any writes to these regions are immediately transferred to the underlying file and made visible to other processes.
- If you don't specify this flag, Linux may buffer writes before transferring them to the file.
- As with shared memory segments, users of memory-mapped regions must establish and follow a protocol to avoid race conditions

Sockets

- A socket is a bidirectional communication mechanism that can be used to communicate with another process on the same machine or with a process running on other machine.

Socket Concepts

- When data is sent through a socket, it is packaged into chunks called packets.
- A socket address struct identifies one end of a socket connection.
- Client is the process initiating the connection, and server is the process waiting to accept connections
- We can read from or write to the socket using read, write calls as with files.
- We can also use functions like send and recv instead of read and write. Send and Recv are specific to socket IO and provide some advanced options.

Socket Concepts (contd)

- Similar to a file, a Socket is represented by file descriptors.
- A port number distinguishes among multiple sockets on the same host.

Eg:

Web servers use port number 80

SMTP servers use port number 25

- For your own server applications, use port numbers greater than 1024

Socket Concepts (contd)

- DNS Names
 - It is easier to remember names than numbers
 - The Domain Name Service (DNS) associates names such as `www.xyz.com` with unique IP numbers
 - To convert human-readable hostnames, either numbers in standard dot notation (such as `10.0.0.1`) or DNS names (such as `www.xyz.com`) into 32-bit IP numbers, you can use `gethostbyname()` function

Socket System Calls

- `socket` — Creates a socket
- `Close` — Destroys a socket
- `connect` — Creates a connection between two sockets
- `bind` — Labels a server socket with an address
- `listen` — Configures a socket to accept connections
- `accept` — Accepts a connection and creates a new socket for the connection

Creating and destroying Sockets

- The `socket()` and `close()` functions create and destroy sockets, respectively
- When you create a socket, you must specify three parameters:
 - Connection type,
 - namespace (local or internet), and
 - Protocol
- Connection type
 - Connection type controls how the socket treats transmitted data and specifies the number of communication partners
 - Connection-oriented type guarantees delivery of all packets in the order they were sent
 - If packets are lost or reordered by problems in the network, the receiver automatically requests their retransmission from the sender

Creating and destroying Sockets (contd)

- Connection type (contd)
 - Connectionless (Datagram) type does not guarantee delivery or arrival order
 - Packets may be lost or reordered in transit due to network errors or other conditions
 - The sender specifies the receiver's address for each individual message
- For specifying connection type parameter, use `SOCK_STREAM` for connection-oriented type, and `SOCK_DGRAM` for Connectionless (Datagram) type

Creating and destroying Sockets (contd)

- **Socket Namespace**
 - A socket namespace specifies how socket addresses are specified
 - For specifying namespace parameter, use
 - PF_LOCAL or PF_UNIX for local namespace, and PF_INET for Internet namespace.
 - Socket addresses in the “local namespace” are ordinary filenames
 - In “Internet namespace”, a socket address is composed of host IP address and a port number.

Creating and destroying Sockets (contd)

- Protocol
 - A protocol specifies how data is transmitted. Eg: TCP/IP, UDP or UNIX local communication protocol, etc.
 - Each protocol is valid for a particular namespace-style combination
 - There is usually one best protocol for each such pair, specifying 0 is recommended
- If call to `socket()` succeeds, it returns a file descriptor for the socket
- For closing the connection, call `close()`

Connect() function

- Client calls connect() function, specifying the address of a server socket to which it wants to connect.
- connect(sock_fd, socket address struct, length of socket address struct)
- Socket address formats differ according to the socket namespace.

Server Socket Calls

- A server does the following :
 - creates a socket,
 - binds an address to its socket,
 - calls `listen()`, that enables connections to the socket,
 - calls `accept()` to accept incoming connections,
 - and then closes the socket
- Data isn't read and written directly via the server socket
- Instead, each time a server program accepts a new connection, Linux creates a separate socket (connected socket) for transferring data
- `bind()` function
 - An address must be bound to the server's socket using `bind`.
 - `bind(sock_fd, socket address struct, length of socket address struct)`

Server Socket Calls (contd)

- After an address is bound to the socket, server must call `listen()`
 - Listen indicates that server is ready to accept connections
 - `listen(sock_fd, num of conn)`
 - The second argument specifies how many pending connections can be queued
 - If the queue is full, additional connections will be rejected
 - Note, this does not limit the total number of connections that a server can handle; it limits just the number of clients attempting to connect that have not yet been accepted

Server Socket Calls (contd)

- A server accepts a connection request from a client by invoking `accept()`
 - `accept(sock_fd, socket address struct, length of socket address struct)`
 - Second argument contains the address of the client
 - The call to `accept` creates a new socket for communicating with the client and returns the corresponding file descriptor
 - The original server socket continues to accept new client connections

LOCAL or UNIX domain sockets

- Sockets connecting processes on the same computer
- They are called local sockets or UNIX-domain sockets
- Their socket addresses are specified by filenames
- The socket's name is specified in struct `sockaddr_un` :
 - set the `sun_family` field to `AF_LOCAL`
 - `sun_path` field specifies the filename
- Use `SUN_LEN` macro to get the length of this socket address struct
- The only permissible protocol value for the local namespace is 0

LOCAL or UNIX domain sockets (contd)

- Because it resides in file system, a local socket is listed as a file

For example,

```
% ls -l /tmp/socket
```

```
srwxrwx--x 1 user group 0 Nov 13 19:18 /tmp/socket
```

- Notice the initial s
- Call unlink to remove a local socket file

Internet-Domain Sockets

- Internet-domain sockets are used to connect processes on different machines connected by a network
- The most common protocols for Internet namespace are TCP/IP
- Internet socket addresses are stored in struct `sockaddr_in` :
 - Set the `sin_family` field to `AF_INET`
 - The `sin_addr` field stores the Internet address of the desired machine as a 32-bit IP address
- Specify port number to differentiate between sockets on the machine.

Review

- Topics Covered
 - What is IPC?
 - Different IPC mechanisms
 - Shared Memory
 - Mapped Memory (mmap)
 - Semaphores
 - Sockets